(https://cognitiveclass.ai)

Convolutional Neural Networks with Keras

In this lab, we will learn how to use the Keras library to build convolutional neural networks. We will also use the popular MNIST dataset and we will compare our results to using a conventional neural network.

Table of Contents

Import Keras and Packages

Let's start by importing the keras libraries and the packages that we would need to build a neural network.

In [1]:

```python
import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.utils import to_categorical
```

```
Using TensorFlow backend.
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorfl
ow/python/framework/dtypes.py:519: FutureWarning: Passing (type, 1) or
'1type' as a synonym of type is deprecated; in a future version of nump
y, it will be understood as (type, (1,)) / '(1,)type'.
  _np_qint8 = np.dtype([("qint8", np.int8, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorfl
ow/python/framework/dtypes.py:520: FutureWarning: Passing (type, 1) or
'1type' as a synonym of type is deprecated; in a future version of nump
y, it will be understood as (type, (1,)) / '(1,)type'.
  _np_quint8 = np.dtype([("quint8", np.uint8, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorfl
ow/python/framework/dtypes.py:521: FutureWarning: Passing (type, 1) or
'1type' as a synonym of type is deprecated; in a future version of nump
y, it will be understood as (type, (1,)) / '(1,)type'.
  _np_qint16 = np.dtype([("qint16", np.int16, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorfl
ow/python/framework/dtypes.py:522: FutureWarning: Passing (type, 1) or
'1type' as a synonym of type is deprecated; in a future version of nump
y, it will be understood as (type, (1,)) / '(1,)type'.
  _np_quint16 = np.dtype([("quint16", np.uint16, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorfl
ow/python/framework/dtypes.py:523: FutureWarning: Passing (type, 1) or
'1type' as a synonym of type is deprecated; in a future version of nump
y, it will be understood as (type, (1,)) / '(1,)type'.
  _np_qint32 = np.dtype([("qint32", np.int32, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorfl
ow/python/framework/dtypes.py:528: FutureWarning: Passing (type, 1) or
'1type' as a synonym of type is deprecated; in a future version of nump
y, it will be understood as (type, (1,)) / '(1,)type'.
  np_resource = np.dtype([("resource", np.ubyte, 1)])
```

When working with convolutional neural networks in particular, we will need additional packages.

In [2]:

```python
from keras.layers.convolutional import Conv2D # to add convolutional layers
from keras.layers.convolutional import MaxPooling2D # to add pooling layers
from keras.layers import Flatten # to flatten data for fully connected layers
```

Convolutional Layer with One set of convolutional and pooling layers

In [3]:

```python
# import data
from keras.datasets import mnist

# load data
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# reshape to be [samples][pixels][width][height]
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32')
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32')
```

```
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz
11493376/11490434 [==============================] - 0s 0us/step
```

Let's normalize the pixel values to be between 0 and 1

In [4]:

```python
X_train = X_train / 255 # normalize training data
X_test = X_test / 255 # normalize test data
```

Next, let's convert the target variable into binary categories

In [5]:

```python
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

num_classes = y_test.shape[1] # number of categories
```

Next, let's define a function that creates our model. Let's start with one set of convolutional and pooling layers.

In [6]:

```python
def convolutional_model():

    # create model
    model = Sequential()
    model.add(Conv2D(16, (5, 5), strides=(1, 1), activation='relu', input_shape=(28
, 28, 1)))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

    model.add(Flatten())
    model.add(Dense(100, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))

    # compile model
    model.compile(optimizer='adam', loss='categorical_crossentropy',  metrics=['acc
uracy'])
    return model
```

Finally, let's call the function to create the model, and then let's train it and evaluate it.

In [7]:

```python
# build the model
model = convolutional_model()

# fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, batch_size
=200, verbose=2)

# evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: {} \n Error: {}".format(scores[1], 100-scores[1]*100))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
 - 136s - loss: 0.2789 - acc: 0.9228 - val_loss: 0.0896 - val_acc: 0.97
38
Epoch 2/10
 - 136s - loss: 0.0784 - acc: 0.9773 - val_loss: 0.0568 - val_acc: 0.98
23
Epoch 3/10
 - 137s - loss: 0.0544 - acc: 0.9837 - val_loss: 0.0460 - val_acc: 0.98
53
Epoch 4/10
 - 137s - loss: 0.0422 - acc: 0.9874 - val_loss: 0.0431 - val_acc: 0.98
62
Epoch 5/10
 - 137s - loss: 0.0346 - acc: 0.9896 - val_loss: 0.0398 - val_acc: 0.98
75
Epoch 6/10
 - 137s - loss: 0.0280 - acc: 0.9913 - val_loss: 0.0394 - val_acc: 0.98
74
Epoch 7/10
 - 137s - loss: 0.0245 - acc: 0.9922 - val_loss: 0.0389 - val_acc: 0.98
75
Epoch 8/10
 - 137s - loss: 0.0200 - acc: 0.9941 - val_loss: 0.0391 - val_acc: 0.98
72
Epoch 9/10
 - 135s - loss: 0.0165 - acc: 0.9951 - val_loss: 0.0406 - val_acc: 0.98
76
Epoch 10/10
 - 135s - loss: 0.0151 - acc: 0.9954 - val_loss: 0.0366 - val_acc: 0.98
80
Accuracy: 0.988
 Error: 1.2000000000000028
```

Convolutional Layer with two sets of convolutional and pooling layers

Let's redefine our convolutional model so that it has two convolutional and pooling layers instead of just one layer of each.

In [8]:

```python
def convolutional_model():

    # create model
    model = Sequential()
    model.add(Conv2D(16, (5, 5), activation='relu', input_shape=(28, 28, 1)))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

    model.add(Conv2D(8, (2, 2), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

    model.add(Flatten())
    model.add(Dense(100, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))

    # Compile model
    model.compile(optimizer='adam', loss='categorical_crossentropy',  metrics=['accuracy'])
    return model
```

Now, let's call the function to create our new convolutional neural network, and then let's train it and evaluate it.

In [ ]:

```python
# build the model
model = convolutional_model()

# fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, batch_size=200, verbose=2)

# evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: {} \n Error: {}".format(scores[1], 100-scores[1]*100))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
 - 141s - loss: 0.4650 - acc: 0.8701 - val_loss: 0.1360 - val_acc: 0.9585
Epoch 2/10
 - 140s - loss: 0.1159 - acc: 0.9647 - val_loss: 0.0865 - val_acc: 0.9743
Epoch 3/10
```

Thank you for completing this lab!

This notebook was created by Alex Aklson (https://www.linkedin.com/in/aklson/). I hope you found this lab interesting and educational. Feel free to contact me if you have any questions!

This notebook is part of a course on **Coursera** called *Introduction to Deep Learning & Neural Networks with Keras*. If you accessed this notebook outside the course, you can take this course online by clicking here (https://cocl.us/DL0101EN_Coursera_Week4_LAB1).