



IBM Developer SKILLS NETWORK

(<https://cognitiveclass.ai>).

Convolutional Neural Networks with Keras

In [10]:

#by Christopher Harrison

In this lab, we will learn how to use the Keras library to build convolutional neural networks. We will also use the popular MNIST dataset and we will compare our results to using a conventional neural network.

Table of Contents

1. [Import Keras and Packages](#) 2. [Convolutional Neural Network with One Convolutional and Pooling Layers](#) 3. [Convolutional Neural Network with Two Convolutional and Pooling Layers](#)

Import Keras and Packages

Let's start by importing the keras libraries and the packages that we would need to build a neural network.

In [11]:

```
import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.utils import to_categorical
```

When working with convolutional neural networks in particular, we will need additional packages.

In [12]:

```
from keras.layers.convolutional import Conv2D # to add convolutional layers
from keras.layers.convolutional import MaxPooling2D # to add pooling layers
from keras.layers import Flatten # to flatten data for fully connected layers
```

Convolutional Layer with One set of convolutional and pooling layers

In [13]:

```
# import data
from keras.datasets import mnist

# load data
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# reshape to be [samples][pixels][width][height]
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32')
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32')
```

Let's normalize the pixel values to be between 0 and 1

In [14]:

```
X_train = X_train / 255 # normalize training data
X_test = X_test / 255 # normalize test data
```

Next, let's convert the target variable into binary categories

In [15]:

```
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

num_classes = y_test.shape[1] # number of categories
```

Next, let's define a function that creates our model. Let's start with one set of convolutional and pooling layers.

In [16]:

```
def convolutional_model():

    # create model
    model = Sequential()
    model.add(Conv2D(16, (5, 5), strides=(1, 1), activation='relu', input_shape=(28, 28, 1)))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

    model.add(Flatten())
    model.add(Dense(100, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))

    # compile model
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
    return model
```

Finally, let's call the function to create the model, and then let's train it and evaluate it.

In [17]:

```
# build the model
model = convolutional_model()

# fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, batch_size=200, verbose=2)

# evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: {} \n Error: {}".format(scores[1], 100-scores[1]*100))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

- 138s - loss: 0.3064 - acc: 0.9131 - val_loss: 0.1075 - val_acc: 0.9690

Epoch 2/10

- 137s - loss: 0.0938 - acc: 0.9731 - val_loss: 0.0650 - val_acc: 0.9804

Epoch 3/10

- 137s - loss: 0.0623 - acc: 0.9816 - val_loss: 0.0513 - val_acc: 0.9843

Epoch 4/10

- 137s - loss: 0.0467 - acc: 0.9860 - val_loss: 0.0437 - val_acc: 0.9865

Epoch 5/10

- 137s - loss: 0.0373 - acc: 0.9888 - val_loss: 0.0424 - val_acc: 0.9865

Epoch 6/10

- 137s - loss: 0.0312 - acc: 0.9901 - val_loss: 0.0373 - val_acc: 0.9879

Epoch 7/10

- 137s - loss: 0.0249 - acc: 0.9926 - val_loss: 0.0383 - val_acc: 0.9886

Epoch 8/10

- 136s - loss: 0.0210 - acc: 0.9939 - val_loss: 0.0388 - val_acc: 0.9870

Epoch 9/10

- 136s - loss: 0.0184 - acc: 0.9944 - val_loss: 0.0401 - val_acc: 0.9866

Epoch 10/10

- 135s - loss: 0.0134 - acc: 0.9964 - val_loss: 0.0369 - val_acc: 0.9886

Accuracy: 0.9886

Error: 1.1400000000000006

Convolutional Layer with two sets of convolutional and pooling layers

Let's redefine our convolutional model so that it has two convolutional and pooling layers instead of just one layer of each.

In [18]:

```
def convolutional_model():  
  
    # create model  
    model = Sequential()  
    model.add(Conv2D(16, (5, 5), activation='relu', input_shape=(28, 28, 1)))  
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))  
  
    model.add(Conv2D(8, (2, 2), activation='relu'))  
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))  
  
    model.add(Flatten())  
    model.add(Dense(100, activation='relu'))  
    model.add(Dense(num_classes, activation='softmax'))  
  
    # Compile model  
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
    return model
```

Now, let's call the function to create our new convolutional neural network, and then let's train it and evaluate it.

In [19]:

```
# build the model
model = convolutional_model()

# fit the model
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10, batch_size=200, verbose=2)

# evaluate the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: {} \n Error: {}".format(scores[1], 100-scores[1]*100))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

- 145s - loss: 0.4829 - acc: 0.8664 - val_loss: 0.1505 - val_acc: 0.95

38

Epoch 2/10

- 143s - loss: 0.1293 - acc: 0.9615 - val_loss: 0.0976 - val_acc: 0.97

12

Epoch 3/10

- 146s - loss: 0.0917 - acc: 0.9729 - val_loss: 0.0785 - val_acc: 0.97

46

Epoch 4/10

- 145s - loss: 0.0735 - acc: 0.9782 - val_loss: 0.0698 - val_acc: 0.97

75

Epoch 5/10

- 146s - loss: 0.0628 - acc: 0.9809 - val_loss: 0.0606 - val_acc: 0.98

12

Epoch 6/10

- 145s - loss: 0.0555 - acc: 0.9829 - val_loss: 0.0560 - val_acc: 0.98

30

Epoch 7/10

- 146s - loss: 0.0491 - acc: 0.9850 - val_loss: 0.0500 - val_acc: 0.98

38

Epoch 8/10

- 145s - loss: 0.0454 - acc: 0.9860 - val_loss: 0.0429 - val_acc: 0.98

63

Epoch 9/10

- 145s - loss: 0.0409 - acc: 0.9874 - val_loss: 0.0411 - val_acc: 0.98

71

Epoch 10/10

- 145s - loss: 0.0358 - acc: 0.9891 - val_loss: 0.0432 - val_acc: 0.98

64

Accuracy: 0.9864

Error: 1.3599999999999994

Thank you for completing this lab!

This notebook was created by [Alex Aklson \(https://www.linkedin.com/in/aklson/\)](https://www.linkedin.com/in/aklson/). I hope you found this lab interesting and educational. Feel free to contact me if you have any questions!

This notebook is part of a course on **Coursera** called *Introduction to Deep Learning & Neural Networks with Keras*. If you accessed this notebook outside the course, you can take this course online by clicking [here](https://cocl.us/DL0101EN_Coursera_Week4_LAB1) (https://cocl.us/DL0101EN_Coursera_Week4_LAB1).

Copyright © 2019 [IBM Developer Skills Network](https://cognitiveclass.ai/?utm_source=bducopyrightlink&utm_medium=dswb&utm_campaign=bdu) (https://cognitiveclass.ai/?utm_source=bducopyrightlink&utm_medium=dswb&utm_campaign=bdu). This notebook and its source code are released under the terms of the [MIT License](https://bigdatauniversity.com/mit-license/) (<https://bigdatauniversity.com/mit-license/>).